

Chunking Strategy Deep Dive

Retrieval-Augmented Generation (RAG) relies heavily on high-quality text chunking.

Standard chunking methods often ignore the structural elements of a document, such as headers,

tables, and lists. This can lead to chunks that are semantically incomplete or confusing to the LLM during the synthesis stage.

Table 1: Chunking Comparison

Strategy	Benefit	Drawback
Fixed-Length	Simple, Fast	Breaks context
Recursive	Structural	Heuristic-based
Semantic	Contextual	Expensive
Layout-Aware	High-Fidelity	Requires OCR/Parsing